

Class 10: Halloween Mini Project

Libby (pid: A69047570)

Libraries

```
library(readr)
library(ggplot2)
library(tidyverse)
```

Importing candy data

```
candy <- read.csv("candy-data.txt", row.names = 1)
head(candy)
```

	chocolate	fruity	caramel	peanut	almond	nougat	crisped	rice	wafer
100 Grand	1	0	1		0	0			1
3 Musketeers	1	0	0		0	1			0
One dime	0	0	0		0	0			0
One quarter	0	0	0		0	0			0
Air Heads	0	1	0		0	0			0
Almond Joy	1	0	0		1	0			0
	hard	bar	pluribus	sugarpercent	pricepercent	winpercent			
100 Grand	0	1	0	0.732	0.860	66.97173			
3 Musketeers	0	1	0	0.604	0.511	67.60294			
One dime	0	0	0	0.011	0.116	32.26109			
One quarter	0	0	0	0.011	0.511	46.11650			
Air Heads	0	0	0	0.906	0.511	52.34146			
Almond Joy	0	1	0	0.465	0.767	50.34755			

```
flextable::flextable(head(candy))
```

chocolate	fruity	caramel	peanut	almond	nougat	crispedrice	wafer	hard	bar	pluribus	s
1	0	1	0	0	0	1	0	0	1	0	0
1	0	0	0	0	1	0	0	0	1	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	1	0	0	0	0	0	0	0	0	0	0
1	0	0	1	0	0	0	0	0	1	0	0

q1: how many different candy types are in this dataset?

```
nrow(candy)
```

[1] 85

q2: how many fruity candy types are in the dataset

```
sum(candy$fruity)
```

[1] 38

Section 2: what is your favorite candy?

q3. What is your favorite candy in the dataset and what is it's winpercent value?

```
candy["Laffy Taffy", ]$winpercent
```

[1] 41.38956

Q4. What is the winpercent value for "Kit Kat"?

```
candy["Kit Kat", ]$winpercent
```

[1] 76.7686

Q5. What is the winpercent value for "Tootsie Roll Snack Bars"?

```
candy["Tootsie Roll Snack Bars", ]$winpercent
```

```
[1] 49.6535
```

```
library("skimr")
# skim() function in the skimr package that can help give you a quick overview of a given dataset
skim(candy)
```

Table 2: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmond	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

winpercent looks like it is on a 0-100% scale whereas the other columns are 0-1.


Q7. What do you think a zero and one represent for the candy\$chocolate column?

They are representing conditionals. TRUE is 1 and FALSE is 0.

A good place to start any exploratory analysis is with a histogram. You can do this most easily with the base R function hist(). Alternatively, you can use ggplot() with geom_hist(). Either works well in this case and (as always) its your choice.

Q8. Plot a histogram of winpercent values

```
ggplot(candy, aes(x=winpercent)) +  
  geom_histogram(bins=20) +  
  geom_vline(xintercept = 50, color="red")
```




Q9. Is the distribution of winpercent values symmetrical?

It's not perfectly symmetrical, let's try a density plot. As you can see it is skewed to the left a bit.

```
ggplot(candy, aes(x=winpercent)) +  
  geom_density(bins=20) +  
  geom_vline(xintercept = 50, color="red")
```

```
Warning in geom_density(bins = 20): Ignoring unknown parameters: `bins`
```



Q10. Is the center of the distribution above or below 50%?

Although the mean is above 50 (slightly at 50.3 %) due to outliers, the center of the distribution (median) is below 50% at 47.82%

```
win <- candy$winpercent
summary(win)
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	22.45	39.14	47.83	50.32	59.86	84.18

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

Chocolate candy is ranked higher than fruit candy on average at a win percent of 60.02% for chocolate compared to 44.12% of fruity candy

```
# 1. Find all chocolate candy in the dataset
# 2. Extract their winpercent values
# 3. Find the mean of the values
chocolate <- candy |>
  filter(chocolate == 1) |>
  summary(winpercent)
```

```

# Mean :60.92

fruity <- candy |>
  filter(fruity == 1) |>
  summary(winpercent)
# Mean :44.12

# Alternative:
# 1. Find all chocolate candy in the dataset
choc inds <- as.logical(candy$chocolate)
choc.candy <- candy[choc inds, ]
# 2. Extract their winpercent values
choc.win <- choc.candy$winpercent
# 3. Find the mean of the values
mean(choc.win)

```

[1] 60.92153

```

# 1. Find all fruity candy in the dataset
fruity.inds <- as.logical(candy$fruity)
fruity.candy <- candy[fruity.inds, ]
# 2. Extract their winpercent values
fruity.win <- fruity.candy$winpercent
# 3. Find the mean of the values
mean(fruity.win)

```

[1] 44.11974

Q12. Is this difference statistically significant?

Yes the answer is statistically significant, so chocolate candy is ranked higher than fruity candy, suggested here by the pvalue of 2.871e-08

```
t.test(choc.win, fruity.win)
```

Welch Two Sample t-test

```

data: choc.win and fruity.win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0

```

```

95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
 60.92153 44.11974

```

3. Overall Candy Rankings

Q13. What are the five least liked candy types in this set?

```
# with base R
head(candy[order(candy$winpercent),], n=5)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat	
Nik L Nip	0	1	0		0	0	
Boston Baked Beans	0	0	0		1	0	
Chiclets	0	1	0		0	0	
Super Bubble	0	1	0		0	0	
Jawbusters	0	1	0		0	0	
	crispedrice	wafer	hard	bar	pluribus	sugar	percent
Nik L Nip	0	0	0	1		0.197	0.976
Boston Baked Beans	0	0	0	1		0.313	0.511
Chiclets	0	0	0	1		0.046	0.325
Super Bubble	0	0	0	0		0.162	0.116
Jawbusters	0	1	0	1		0.093	0.511
	winpercent						
Nik L Nip	22.44534						
Boston Baked Beans	23.41782						
Chiclets	24.52499						
Super Bubble	27.30386						
Jawbusters	28.12744						

```
# with dplyr
candy |>
  arrange(winpercent) |>
  head(5)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat
Nik L Nip	0	1	0		0	0
Boston Baked Beans	0	0	0		1	0
Chiclets	0	1	0		0	0

Super Bubble	0	1	0	0	0	
Jawbusters	0	1	0	0	0	
crispedricewafer hard bar pluribus sugarpercent pricepercent						
Nik L Nip	0	0	0	1	0.197	0.976
Boston Baked Beans	0	0	0	1	0.313	0.511
Chiclets	0	0	0	1	0.046	0.325
Super Bubble	0	0	0	0	0.162	0.116
Jawbusters	0	1	0	1	0.093	0.511
winpercent						
Nik L Nip	22.44534					
Boston Baked Beans	23.41782					
Chiclets	24.52499					
Super Bubble	27.30386					
Jawbusters	28.12744					


Q14. What are the top 5 all time favorite candy types out of this set?

```
candy |>
  arrange(-winpercent) |>
  head(5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Reese's Peanut Butter cup	1	0	0	1	0	
Reese's Miniatures	1	0	0	1	0	
Twix	1	0	1	0	0	
Kit Kat	1	0	0	0	0	
Snickers	1	0	1	1	1	
crispedricewafer hard bar pluribus sugarpercent						
Reese's Peanut Butter cup	0	0	0	0	0.720	
Reese's Miniatures	0	0	0	0	0.034	
Twix	1	0	1	0	0.546	
Kit Kat	1	0	1	0	0.313	
Snickers	0	0	1	0	0.546	
pricepercent winpercent						
Reese's Peanut Butter cup	0.651	84.18029				
Reese's Miniatures	0.279	81.86626				
Twix	0.906	81.64291				
Kit Kat	0.511	76.76860				
Snickers	0.651	76.67378				

Q15. Make a first barplot of candy ranking based on winpercent values

```
ggplot(candy) +
  aes(winpercent, rownames(candy)) +
  geom_col()
```



```
x <- c(10,2,5,1)
order(x)
```

```
[1] 4 2 3 1
```


```
ord.ind <- order(candy$winpercent)
head(candy[ord.ind, ], 5)
```

	chocolate	fruity	caramel	peanut	almondy	nougat	
Nik L Nip	0	1	0	0	0	0	
Boston Baked Beans	0	0	0	1	0	0	
Chiclets	0	1	0	0	0	0	
Super Bubble	0	1	0	0	0	0	
Jawbusters	0	1	0	0	0	0	
	crisped rice	wafer	hard	bar	pluribus	sugarpercent	pricepercent
Nik L Nip	0	0	0	1	0	0.197	0.976
Boston Baked Beans	0	0	0	1	0	0.313	0.511

Chiclets	0	0	0	1	0.046	0.325
Super Bubble	0	0	0	0	0.162	0.116
Jawbusters	0	1	0	1	0.093	0.511
winpercent						
Nik L Nip	22.44534					
Boston Baked Beans	23.41782					
Chiclets	24.52499					
Super Bubble	27.30386					
Jawbusters	28.12744					


Q16. This is quite ugly, use the reorder() function to get the bars sorted by winpercent?

```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col()
```



```
my_cols=rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] = "chocolate"
my_cols[as.logical(candy$bar)] = "brown"
my_cols[as.logical(candy$fruity)] = "pink"
```

```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col(fill=my_cols)
```



Q17. What is the worst ranked chocolate candy?

Sixlets are the worst ranked chocolate candy!


```
my_cols <- rep("black", nrow(candy))
candy$chocolate==1
```

```
[1] TRUE TRUE FALSE FALSE FALSE TRUE TRUE FALSE FALSE FALSE TRUE FALSE
[13] FALSE TRUE TRUE
[25] TRUE TRUE FALSE TRUE TRUE FALSE FALSE FALSE TRUE TRUE FALSE TRUE
[37] TRUE TRUE TRUE TRUE TRUE FALSE TRUE TRUE FALSE FALSE FALSE TRUE
[49] FALSE FALSE FALSE TRUE TRUE TRUE TRUE FALSE TRUE FALSE FALSE TRUE
[61] FALSE FALSE TRUE FALSE TRUE TRUE FALSE FALSE FALSE FALSE FALSE FALSE
[73] FALSE FALSE TRUE TRUE TRUE TRUE FALSE TRUE FALSE FALSE FALSE FALSE
[85] TRUE
```

```
my_cols
```

```
[1] "black" "black" "black" "black" "black" "black" "black" "black" "black"  
[10] "black" "black" "black" "black" "black" "black" "black" "black" "black"  
[19] "black" "black" "black" "black" "black" "black" "black" "black" "black"  
[28] "black" "black" "black" "black" "black" "black" "black" "black" "black"  
[37] "black" "black" "black" "black" "black" "black" "black" "black" "black"  
[46] "black" "black" "black" "black" "black" "black" "black" "black" "black"  
[55] "black" "black" "black" "black" "black" "black" "black" "black" "black"  
[64] "black" "black" "black" "black" "black" "black" "black" "black" "black"  
[73] "black" "black" "black" "black" "black" "black" "black" "black" "black"  
[82] "black" "black" "black" "black"
```

```
ggplot(candy) +  
  aes(winpercent, reorder(rownames(candy), winpercent)) +  
  geom_col(fill=my_cols)
```



Q18. What is the best ranked fruity candy?

Starburst is the best ranked fruity candy!

Section 4


Winpercent v PricePercent

Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

The fruity candies offer more bang for your buck, with Starburst seeming to be the highest ranked for best cost, reeses miniatures may contend with it too, but it does have a slightly higher price percent.

```
library(ggrepel)
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols)
```

Warning: ggrepel: 50 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
ord <- order(candy$pricepercent, decreasing = TRUE)
head( candy[ord,c(11,12)], n=5 )
```

	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

Correlation structure

Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?


fruity and chocolate are anti-correlated, also pluribus and bar but less so than fruity and chocolate

```
cij <- cor(candy) # takes all cols in dataset, so an nxn correlation matrix
```

```
library(corrplot)
```

```
corrplot 0.95 loaded
```

```
corrplot(cij)
```



Q23. Similarly, what two variables are most positively correlated?

Aside from the 1:1 match, when items are paired against themselves, chocolate and winpercent are most positively correlated.

6. Principal Component Analysis

The main function in base R for this `prcomp()` remember it has a `scale` parameter you basically always want to use; `scale=TRUE`

```
pca <- prcomp(candy, scale=TRUE)
summary(pca)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317		
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000		

Let's first look at our first main result figure - the "PC Plot" or PC1 v PC2 plot

```
ggplot(pca$x) + # shows where candies lie
  aes(PC1, PC2, label=rownames(pca$x)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols) +
  labs(title = "Barry's candy space map")
```

Warning: ggrepel: 54 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Don't forget about your variable loadings – how the original variables contribute to your new PCs....

Q24. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you?

Fruity, hard, and pluribus are most strongly picked up by PC1 in the positive direction. This makes sense because most fruity candy does seem to be hard, and because it is cheaper comes in multiples more than chocolate candy.

```
ggplot(pca$rotation) +  
  aes(PC1, rownames(pca$rotation)) +  
  geom_col()
```

